FAIRLENS

Sign In Get Started

Audit Session

ID: 35CE3D

- OVERVIEW
- METRICS
- FEATURE IMPORTANCE
- GEMINI REPORT
- REMEDIATION

RISK ASSESSMENT

LOW SEVERITY

0 of 0 metrics failed.

Audit Summary

0 of 0 core fairness metrics failed. race and age_cat and sex are significant predictors of outcome.

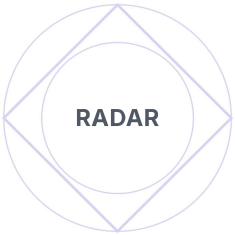

FAIRNESS METRICS APR 26, 2026

Outcome Rate by Protected Group

Outcome Rate by Group

Percentage of positive outcomes per protected group

Export Comprehensive Report

FAIRLENS

Audit Session

ID: 35CE3D

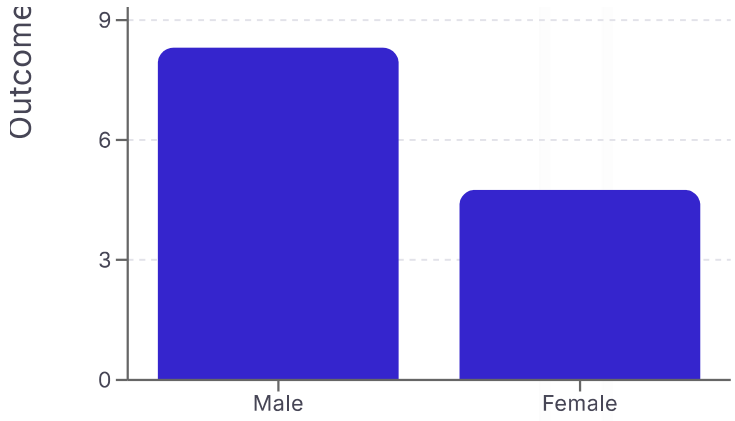

Positive outcome rates per group in the sex protected attribute

What's Driving the Model?

No SHAP data available.

GEMINI REPORT

AI-Generated Audit Findings

SUMMARY

This model is being checked for fairness, but the report we received is missing all the critical information needed to make a judgment. It's like asking a doctor for a diagnosis without providing any test results. However, based on the task—predicting future violence—and the groups involved, there is a very high risk of serious

FAIRLENS

Audit Session

ID: 35CE3D

particularly African-American individuals, as being at high risk of re-offending when they are not. This type of error, a 'false positive,' can have damaging real-world consequences. This bias often happens because the historical data used to 'teach' the model is itself biased, reflecting historical inequalities in the justice system. The model may not use 'race' as a direct input, but it can learn to associate it with other data points like prior arrest records or zip codes. Before this model can be considered safe to use, a complete audit with all the necessary data and fairness test results must be performed. Based on those results, we have clear steps to fix the data and the model to ensure it makes decisions more equitably for all groups.

WHO IS AFFECTED

African-American | Male | Less than 25

ROOT CAUSE

A definitive root cause analysis is impossible due to the absence of SHAP feature importance data. However, potential causes for bias in such models typically include: 1) The model learning proxies for race from correlated features (e.g., prior arrest history, socioeconomic factors). 2) Significant data quality issues, including extensive missing data (e.g., 17,496 missing 'violent_recid' targets)

FAIRLENS

Audit Session

ID: 35CE3D

Historical bias in the training data, where certain groups are disproportionately represented in arrest records, which the model learns and perpetuates.

SUGGESTED:

How can I mitigate race proxy bias?

Which group is most affected by this model?

What are the top remediation steps I should take?

Ask Gemini about this au

STRATEGIC REMEDIATION PROTOCOL

01 Halt model deployment and conduct a full data audit. Investigate and rectify the cause of extensive missing data, especially for the target variable 'is_violent_recid', without which a valid model cannot be trained or audited.

02 Re-run the model audit to generate essential fairness metrics (e.g., False Positive Rate Parity, Equal Opportunity

FAIRLENS

Audit Session

ID: 35CE3D

analysis.

03 Implement targeted bias mitigation techniques based on complete audit findings. This may include pre-processing (e.g., re-weighting data for affected groups), in-processing (e.g., adversarial debiasing), or post-processing (e.g., group-specific decision thresholds).

04 Scrutinize input features for their potential as proxies for protected attributes. Remove or transform features that are highly correlated with race or sex but have low predictive value for the actual outcome.

05 Establish an ongoing monitoring framework to track model performance and fairness metrics over time, ensuring that biases do not re-emerge as data distributions shift.

FAIRLENS

Audit Session

ID: 35CE3D